

GPU Memory Model

Adapted from:

Aaron Lefohn

University of California, Davis

With updates from slides by

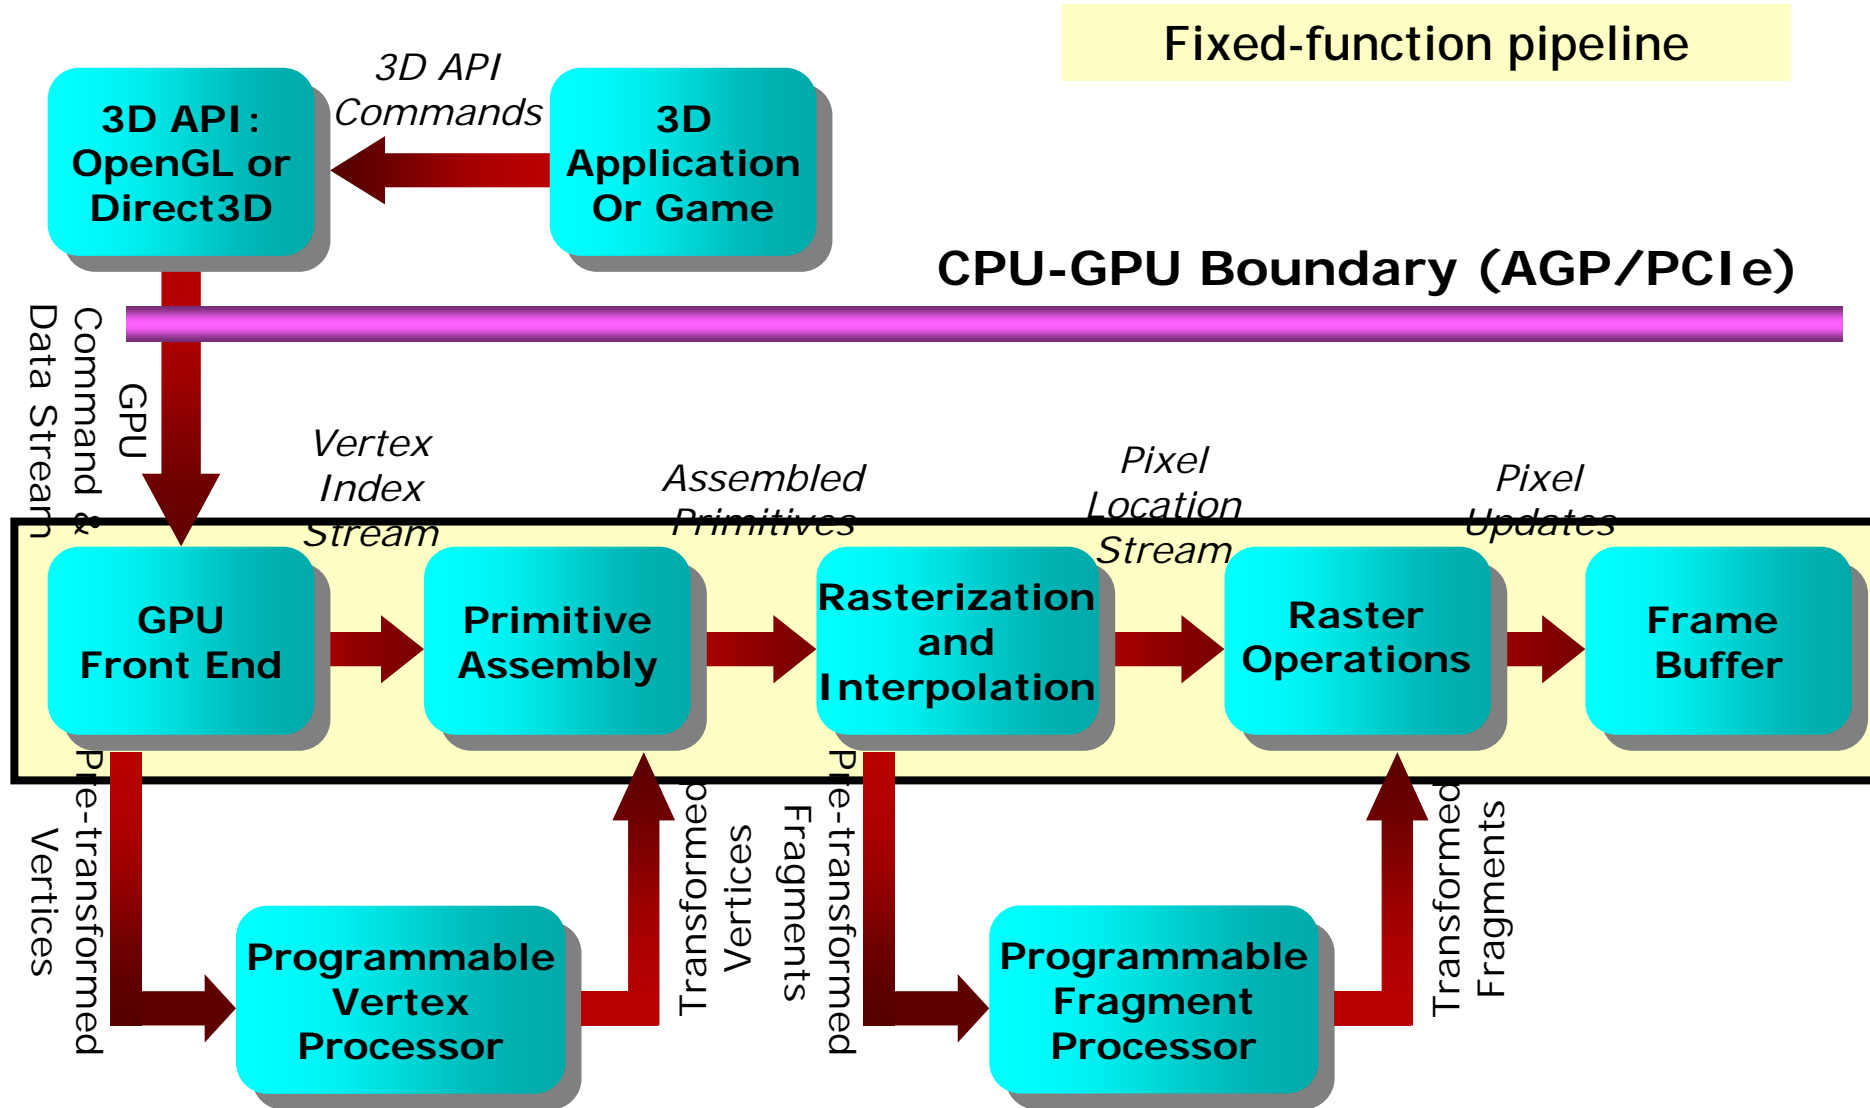
Suresh Venkatasubramanian,

University of Pennsylvania

Updates performed by Gary J. Katz,

University of Pennsylvania

Review




Overview

- GPU Memory Model
- GPU Data Structure Basics
- Introduction to Framebuffer Objects
- Fragment Pipeline
- Vertex Pipeline

Memory Hierarchy

- CPU and GPU Memory Hierarchy



CPU Memory Model


- At any program point
 - Allocate/free local or global memory
 - Random memory access
 - Registers
 - Read/write
 - Local memory
 - Read/write to stack
 - Global memory
 - Read/write to heap
 - Disk
 - Read/write to disk

GPU Memory Model

- Much more restricted memory access
 - Allocate/free memory only before computation
 - Limited memory access during computation (kernel)
 - Registers
 - Read/write
 - Local memory
 - Does not exist
 - Global memory
 - Read-only during computation
 - Write-only at end of computation (pre-computed address)
 - Disk access
 - Does not exist

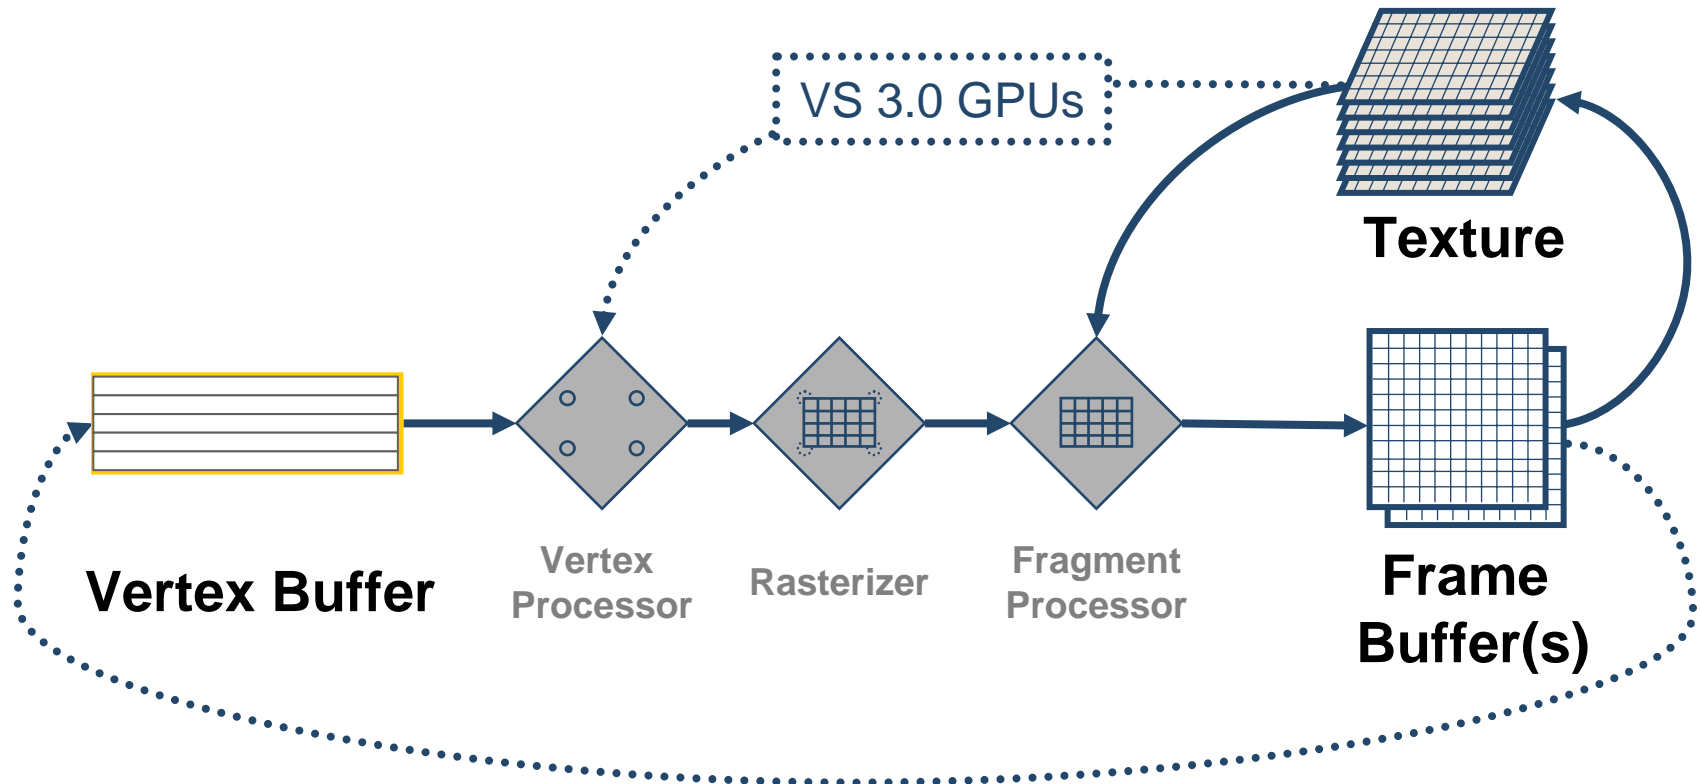
GPU Memory Model

- Where is GPU Data Stored?
 - Vertex buffer
 - Frame buffer
 - Texture



Vertex Buffers

- GPU memory for vertex data
- Vertex data required to initiate render pass



Vertex Buffers

- Supported Operations
 - CPU interface
 - Allocate
 - Free
 - Copy CPU → GPU
 - Copy GPU → GPU (Render-to-vertex-array)
 - Bind for read-only vertex stream access
 - GPU interface
 - Stream read (vertex program only)

Vertex Buffers

- Limitations

- CPU

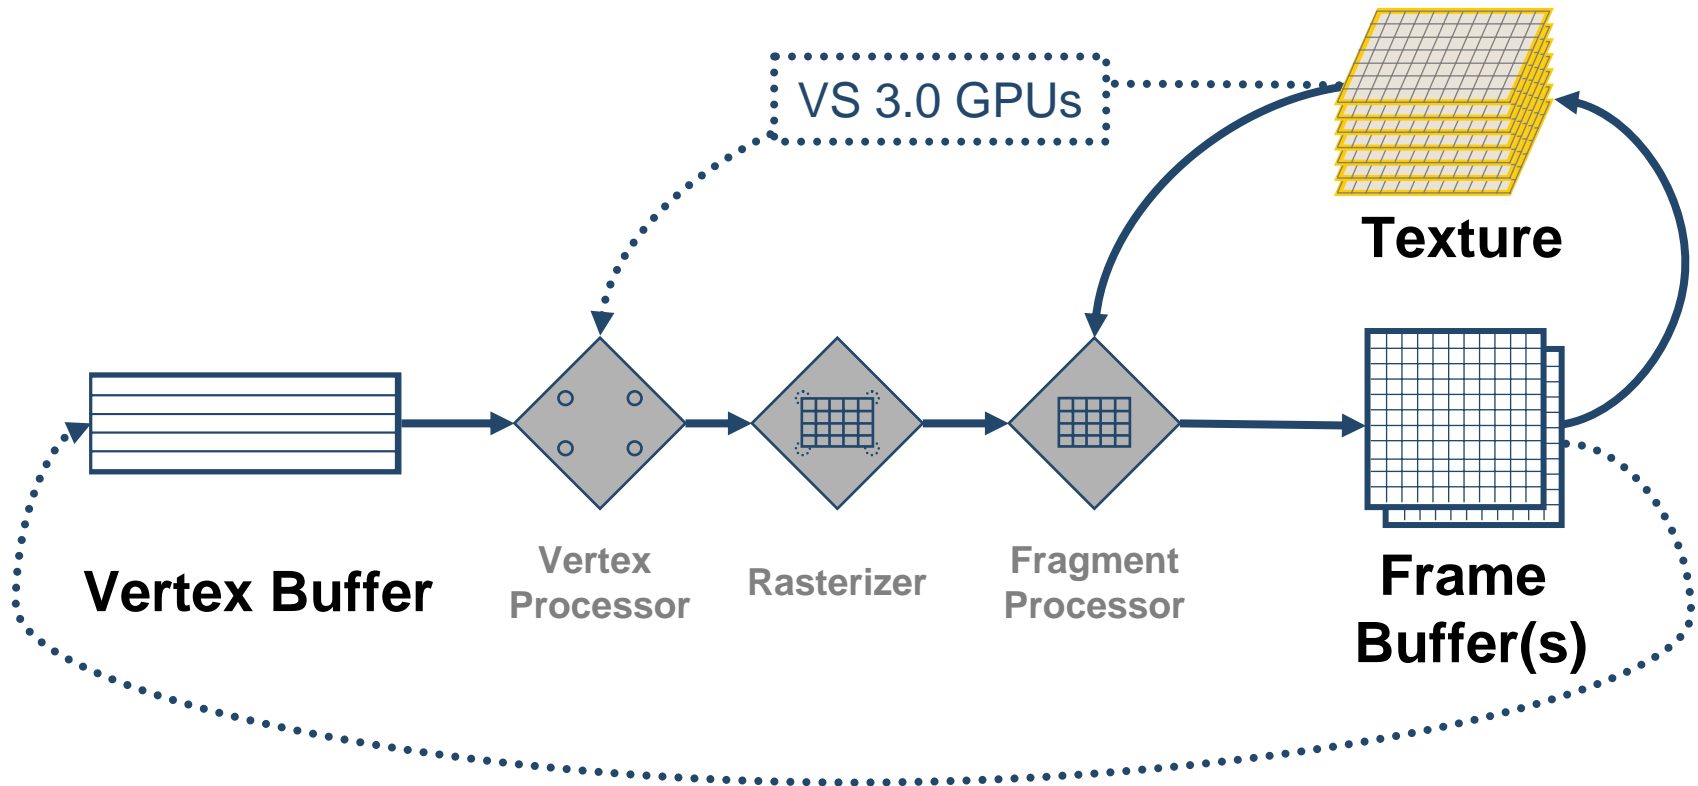
- No copy GPU → CPU
 - No bind for read-only random access

- GPU

- No random-access reads
 - No access from fragment programs

Textures

- Random-access GPU memory



Textures

- Supported Operations
 - CPU interface
 - Allocate
 - Free
 - Copy CPU → GPU
 - Copy GPU → CPU
 - Copy GPU → GPU (Render-to-texture)
 - Bind for read-only random access (vertex or fragment)
 - GPU interface
 - Random read

Textures


- Limitations
 - No bind for vertex stream access

Next ...

- GPU Data Structure Basics

GPU Arrays

- Large 1D Arrays
 - Current GPUs limit 1D array sizes
 - Pack into 2D memory
 - 1D-to-2D address translation



GPU Arrays

- 3D Arrays



- Problem

- GPUs do not have 3D frame buffers

- Solutions

1. Stack of 2D slices

2. Multiple slices per 2D buffer



GPU Arrays

- **Problems With 3D Arrays**
 - Cannot read stack of 2D slices as 3D texture
 - Must know which slices are needed in advance
- **Solutions**
 - Flat 3D textures
 - Need packing functions for coordinate transfer

GPU Arrays

- Higher Dimensional Arrays
 - Pack into 2D buffers
 - N-D to 2D address translation
 - Same problems as 3D arrays if data does not fit in a single 2D texture

Sparse/Adaptive Data Structures

- Why?

- Reduce memory pressure
- Reduce computational workload


- Examples

- Sparse matrices
 - Krueger et al., Siggraph 2003
 - Bolz et al., Siggraph 2003
- Deformable implicit surfaces (sparse volumes/PDEs)
 - Lefohn et al., IEEE Visualization 2003 / TVCG 2004
- Adaptive radiosity solution (Coombe et al.)



Sparse/Adaptive Data Structures

- Basic Idea
 - Pack “active” data elements into GPU memory



GPU Data Structures


- **Conclusions**

- Fundamental GPU memory primitive is a fixed-size 2D array
- GPGPU needs more general memory model
- Building and modifying complex GPU-based data structures is an open research topic...

Framebuffer Objects

- What is an FBO?

- A struct that holds pointers to memory objects
- Each bound memory object can be a framebuffer rendering surface
- Platform-independent



Framebuffer Objects

- Which memory can be bound to an FBO?
 - Textures
 - Renderbuffers
 - Depth, stencil, color
 - Traditional write-only framebuffer

The fragment pipeline

Input: Fragment ■


Attributes

Color	R	G	B	A
Position	X	Y	Z	W
Texture coordinates	X	Y	[Z]	-
Texture coordinates	X	Y	[Z]	-
...				

32 bits = float

16 bits = half

Input: Texture Image



- Each element of texture is 4D vector
- Textures can be “square” or rectangular (power-of-two or not)

The fragment pipeline

Input: Uniform parameters

- Can be passed to a fragment program like normal parameters
- set in advance before the fragment program executes

Example:

A counter that tracks which pass the algorithm is in.

Input: Constant parameters

- Fixed inside program
- E.g. float4 v = (1.0, 1.0, 1.0, 1.0)

Examples:

3.14159..

Size of compute window

The fragment pipeline

Math ops: USE THEM !

- `cos(x)/log2(x)/pow(x,y)`
- `dot(a,b)`
- `mul(v, M)`
- `sqrt(x)`
- `cross(u, v)`

Using built-in ops is more efficient than writing your own

Swizzling/masking: an easy way to move data around.

```
v1 = (4,-2,5,3); // Initialize  
v2 = v1.yx;      // v2 = (-2,4)  
s = v1.w;        // s = 3  
v3 = s.rrr;      // v3 = (3,3,3)
```

Write masking:

```
v4 = (1,5,3,2);  
v4.ar = v2;      // v4=(4,5,4,-2)
```

The fragment pipeline


```
float4 v = tex2D(IMG, float2(x,y))
```

Texture access is like an array lookup.

The value in v can be used to perform another lookup!

This is called a **dependent read**

Texture reads (and dependent reads) are expensive resources, and are limited in different GPUs. Use them wisely !



The fragment pipeline

Control flow:

- (<test>)?a:b operator.
- if-then-else conditional
 - [nv3x] Both branches are executed, and the condition code is used to decide which value is used to write the output register.
 - [nv40] True conditionals
- for-loops and do-while
 - [nv3x] limited to what can be unrolled (i.e no variable loop limits)
 - [nv40] True looping.

The fragment pipeline

Fragment programs use **call-by-result**

```
out float4 result : COLOR
// Do computation
result = <final answer>
```

Notes:

- Only output color can be modified
- Textures cannot be written
- Setting different values in different channels of result can be useful for debugging

The Vertex Pipeline

Input: vertices

- position, color, texture coords.

Input: uniform and constant parameters.

- Matrices can be passed to a vertex program.
- Lighting/material parameters can also be passed.

The Vertex Pipeline

Operations:

- Math/swizzle ops
- Matrix operators
- Flow control (as before)

[nv3x] No access to textures.

Output:

- Modified vertices (position, color)
- Vertex data transmitted to primitive assembly.

Vertex programs are useful

- We can replace the entire geometry transformation portion of the fixed-function pipeline.
- Vertex programs used to change vertex coordinates (move objects around)
- There are many fewer vertices than fragments: shifting operations to vertex programs improves overall pipeline performance.
- Much of shader processing happens at vertex level.
- We have access to original scene geometry.

Vertex programs are not useful

- Fragment programs allow us to exploit full parallelism of GPU pipeline (“a processor at every pixel”).

Rule of thumb:

If computation requires intensive calculation,
it should probably be in the fragment processor.

If it requires more geometric/graphic computing,
it should be in the vertex processor.

Conclusions

- **GPU Memory Model Evolving**
 - Writable GPU memory forms loop-back in an otherwise feed-forward pipeline
 - Memory model will continue to evolve as GPUs become more general data-parallel processors
- **Data Structures**
 - Basic memory primitive is limited-size, 2D texture
 - Use address translation to fit all array dimensions into 2D