


Open CL



Jens Ogniewski
Information Coding Group

Questions

Which types of synchronization does OpenCL offer? Which not?

What is memory coalescing? How can it be achieved?

What characteristics must a function / an algorithm have to be suitable as an OpenCL kernel?

Outline

Introduction

Language Overview

Open CL for NVIDIA GPUs

Dos and donts

An example


Introduction

- Language aimed for parallel architectures
 - Programmer defines explicitly where and how parallelism occurs
 - Aimed towards SIMD processing
 - Task parallelism?
- Heterogeneous = Hardware independent (mostly)
 - Optimized code still needs intimate knowledge of used architecture (GPU vs Cell vs hyperthreaded CISC)
- Based on C99 (with modifications)
- First initiated by Apple, still subject to changes

Introduction



Introduction



The Language


- Based on C99, but:
 - No function pointers (will come later?)
 - No pointers to pointers in function calls (=> no multidimensional arrays, could use image types instead of)
 - No recursion
 - No arrays with dynamical length
 - No bitfields
- Optional:
 - Pointers with length <32 bit
 - Writing support for 3D images
 - Double and half types
 - Atomic functions

The Language

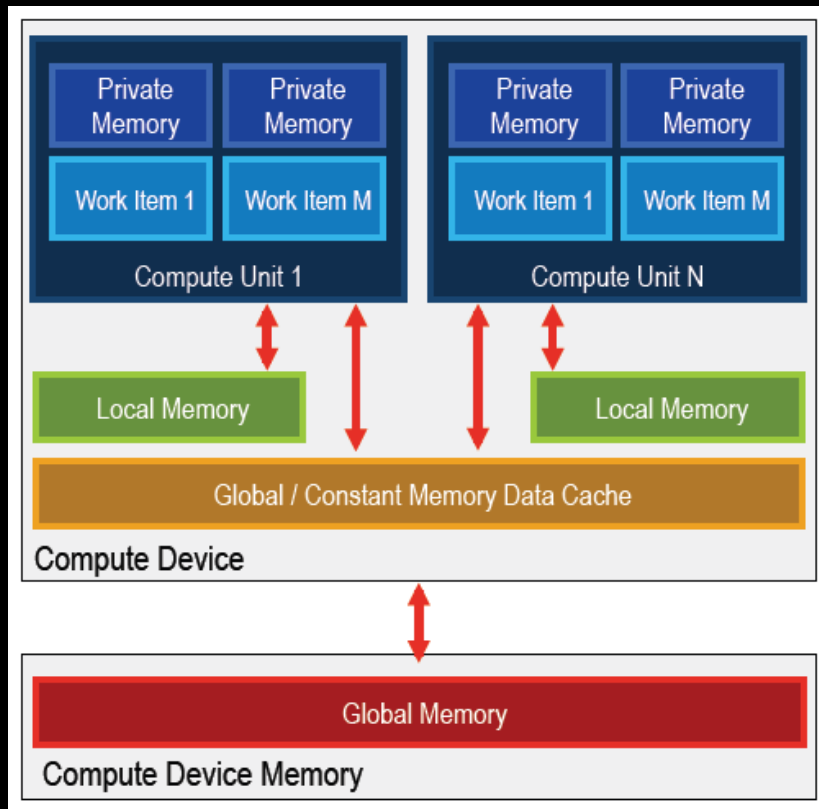
- But instead:
 - Integrated functions for reading / writing 2D images and reading 3D images
 - Converting functions incl. explicit rounding and saturation
 - math.h, all functions with different precisions
 - Vector support (2-, 3- and 4-dimensional)
- Available primitive datatypes:
 - Bool, char, int, long, float, size_t, void, unsigned versions as well
- Mix of OpenCL and OpenGL possible
 - Can share data structures and variables (without copying)
 - API functions available

Platform model

- One host (e.g. PC), one or several compute devices (e.g. graphic card)
 - ▣ Each compute device: one or several compute units (e.g. shader)
 - Each compute device: one or several processing elements



Memory model model

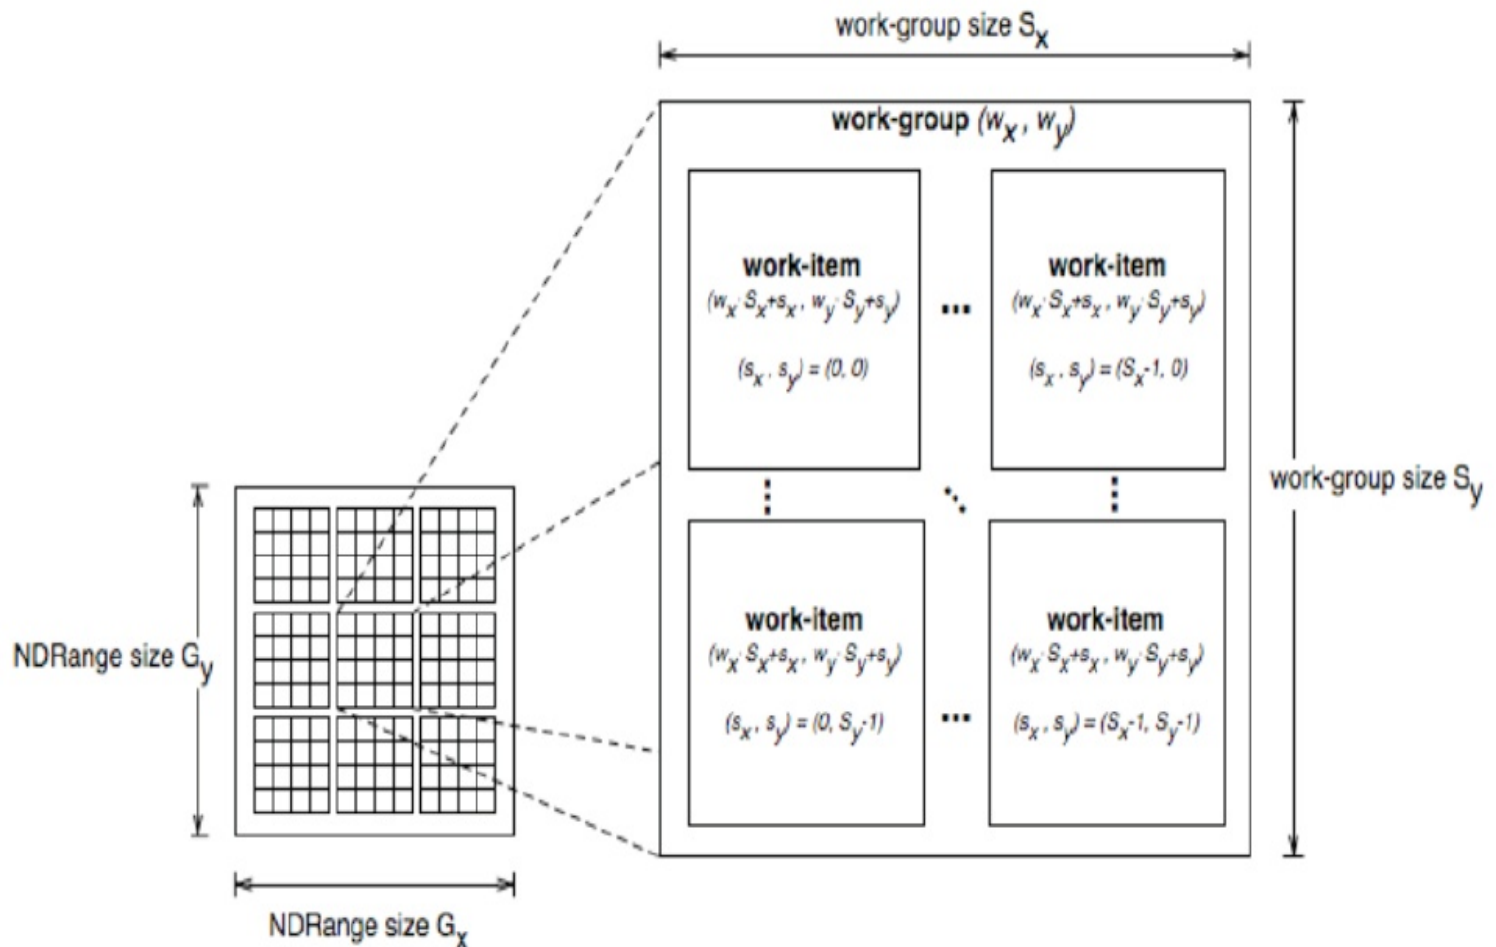


- **Private memory**
 - Only accessible by one processing element (e.g. register)
- **Local memory**
 - Only accessible by one compute unit
 - Copying betw. global/constant memory and local memory has to be done explicitly!
- **Global / Constant memory**
 - Accessible by one compute device
 - Also used by compiler for variables that don't fit into private memory

Execution model

- Kernel
 - (Short) function which will be executed in parallel
 - How many is determined by the global dimensions
 - Each execution is called a work item
- Several work items share one compute unit
 - Called a work group
 - How many is determined by the local dimensions
- Local and global dimensions are given by the programmer
 - Should be chosen carefully depending on application, algorithm and hardware used

Execution model



Execution model

- How does that work?
 - Functions which return position of work-item in workgroup or globally or position of workgroup in global dimensions
 - Can be used to calculate offset

Synchronization

- Memory
 - Explicit data movement between host and compute device memory
 - Explicit data movement between global/constant and local memory on the device
 - Global/constant memory can be “linked” to host memory during creation
- Task synchronization: only inside a workgroup
 - Using `barrier(CLK_LOCAL_MEM_FENCE)`
 - Execution continues only after all items in the workgroup passed this command and wrote back all changes from private memory to local or global memory
 - All work items have to pass the barrier => not allowed inside branches!
- No synchronization between workgroups!

Synchronization

- Between host and compute device(s):
 - Using queues
 - Available for tasks (`clEnqueueNDRangeKernel`)
 - Memory (e.g. `clEnqueueReadBuffer`)
 - And events (e.g. `clWaitforEvents`)
- Queues can be in-order or out-of-order
- Several queues in parallel possible, programmer has to take care of synchronization however

OpenCL for NVIDIA GPUs

- Different terms (CUDA legacy)
 - Compute unit = multiprocessor
 - Work item = thread
 - Work group = thread block, sometimes also warp
 - Carefully: warp often used with a constant size (e.g. 1 warp = 32 threads)
 - And CUDA local memory \neq OpenCL local memory (= CUDA shared memory)
 - CUDA local memory: in global memory (variables that don't fit in the register file)
 - Global memory, constant memory: same in CUDA and OpenCL

OpenCL for NVIDIA GPUs

- Resident work item
 - Has reserved private memory in multiprocessor
 - Does not need to be active, but can become active anytime
- Active work item
 - Executes instruction
- Thread scheduler
 - Can swap out / in resident work items without any overhead
 - Tries to swap out work items waiting for memory access and run other resident work items instead
 - Tries to coalesce memory access inside a workgroup

OpenCL for NVIDIA GPUs

- Example architecture
 - ❑ 8192 registers / compute unit
 - ❑ 16 kb local memory / compute unit
 - ❑ 64 kb constant memory (varying global memory size)
 - ❑ Max. 16 kbytes private memory / working item
 - ❑ Local memory access time: 24 cycles
 - ❑ Global memory access time: 400-600 cycles
 - ❑ Kernel size limit: 2 million PTX instructions
 - ❑ 8 processing elements / compute unit
 - ❑ Max. 768 resident work items per compute unit
 - ❑ Max. 512 work items / work group
 - ❑ Support for atomic functions on 32byte words in global memory
- NVIDIA NVS 290 (your machine): 2 compute units

OpenCL for NVIDIA GPUs

- Example architecture
 - ❑ 8192 registers / compute unit
 - ❑ 16 kb local memory / compute unit
 - ❑ 64 kb constant memory (varying global memory size)
 - ❑ Max. 16 kbytes private memory / working item
 - ❑ Local memory access time: 24 cycles
 - ❑ Global memory access time: 400-600 cycles
 - ❑ Kernel size limit: 2 million PTX instructions
 - ❑ 8 processing elements / compute unit
 - ❑ Max. 768 resident work items per compute unit
 - ❑ Max. 512 work items / work group
 - ❑ Support for atomic functions on 32byte words in global memory
- NVIDIA GTS 250 (Olympen): 16 compute units

Dos & Donts

- When to use GPGPU?
 - Parallel algorithm
 - Instruction mix: little memory access, much computation
 - Data, not task parallel

Or if your algorithm can be rewritten to fulfill those criteria without introducing much overhead (little is ok)!

Dos & Donts

- Avoid:
 - Branches
 - Double precision (at least for now)
 - Memory access: recomputation might lead to faster results
 - Memory bank conflicts (betw. work groups)
 - Private memory in global memory


Dos & Donts

- What else?
 - Use vector intrinsics: to make sure that SIMDs are used correctly and most efficiently
 - Coalescing memory: enforce memory access aligned by 16 words for 16 work items (64 and 128 bytes optimal)
 - Prefer constant to global memory, since it is cached
 - Use local memory as buffer for global memory
 - 16 or 32 work items per work group optimal
 - try to have more than 192 resident work items / compute unit to hide memory accesses
 - Use auto synchronization to avoid barriers
 - Try to reuse kernels as much as possible => compilation expensive
 - Watchdog timer: might disrupt computation

An example



An example



An example: c function

```
for (i=0;i<height*width*4;i++) {  
    int tmp = (*(imageData+i))*(*(lightData+i))/128;  
    if (tmp>255) tmp = 255;  
    *(outputImage+i) = (unsigned char)tmp;  
}
```

An example: kernel function

```
imageShader( __global unsigned char *i_image,  
             __global unsigned char *i_lightmap,  
             __global unsigned char *o_image)  
{  
    int i = get_global_id(0);  
  
    o_image[i] = convert_uchar_sat(i_image[i]  
                                   * i_lightmap[i]/128);  
}
```


An example: init function (1/2)

```
cl_context cxGPUContext;  
cl_command_queue commandQueue;  
cl_kernel gpgpuKernel;  
  
int init_OpenCL(const char* source_path) {  
    cl_int ciErrNum = CL_SUCCESS;  
    char *source;  
    cl_device_id device;  
    size_t program_length;  
  
    cxGPUContext = clCreateContextFromType(0,  
        CL_DEVICE_TYPE_GPU, NULL, NULL, &ciErrNum);
```

An example: init function (2/2)

```
device = oclGetDev(cxGPUContext, 0);
commandQueue = clCreateCommandQueue(cxGPUContext,
    device, 0, &ciErrNum);
source = oclLoadProgSource(source_path,
    "// No header needed!\n", &program_length);
cl_program cpProgram =
    clCreateProgramWithSource(cxGPUContext, 1,
        (const char **)&source, &program_length,
        &ciErrNum);
ciErrNum = clBuildProgram(cpProgram, 0, NULL, NULL,
    NULL, NULL);
gpgpuKernel = clCreateKernel(cpProgram, "imageShader",
    &ciErrNum);
free(source);
}
```

An example: kernel call (1/2)

```
i_image = clCreateBuffer(cxGPUContext,
    CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR, width *
    height * 4 * sizeof(cl_char), image, &ciErrNum);
i_lightmap = clCreateBuffer(cxGPUContext,
    CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR, width *
    height * 4 * sizeof(cl_char), lightmap, &ciErrNum);
o_image = clCreateBuffer(cxGPUContext,
    CL_MEM_READ_WRITE, width * height * 4 *
    sizeof(cl_char), NULL, &ciErrNum);
clSetKernelArg(gpgpuKernel, 0, sizeof(cl_mem),
    (void *) &i_image);
clSetKernelArg(gpgpuKernel, 1, sizeof(cl_mem),
    (void *) &i_lightmap);
clSetKernelArg(gpgpuKernel, 2, sizeof(cl_mem),
    (void *) &o_image);
```

An example: kernel call (2/2)

```
localWorkSize[0] = 1;
globalWorkSize[0] = width * height * 4;
clErrNum = clEnqueueNDRangeKernel(commandQueue,
    gpgpuKernel, 1, NULL, globalWorkSize,
    localWorkSize, 0, NULL, NULL);
clEnqueueReadBuffer(commandQueue, o_image, CL_TRUE, 0,
    width * height * 4 * sizeof(cl_char), result, 0,
    NULL, NULL);

clReleaseMemObject(i_image);
clReleaseMemObject(i_lightmap);
clReleaseMemObject(o_image);
```

Results



- CPU: 18783 μ s
- GPU: 32946 μ s

Results

```
localWorkSize[0] = 16;  
globalWorkSize[0] = width * height * 4;  
ciErrNum = clEnqueueNDRangeKernel(commandQueue,  
    gpgpuKernel, 1, NULL, globalWorkSize,  
    localWorkSize, 0, NULL, NULL);
```

- CPU: 18783 us
- GPU, optimized: 15776 us

Why? Coalesced memory!

Further readings

- Open CL at Khronos

http://www.khronos.org/developers/library/overview/opengl_overview.pdf

<http://www.khronos.org/registry/cl/specs/opengl-1.0.48.pdf>

<http://www.khronos.org/opengl/sdk/1.0/docs/man/xhtml/>

Further readings

- NVIDIA and OpenCL

http://www.nvidia.com/content/cudazone/download/OpenCL/NVIDIA_OpenCL_ProgrammingGuide.pdf

http://www.nvidia.com/content/cudazone/CUDABrowser/downloads/papers/NVIDIA_OpenCL_BestPracticesGuide.pdf

Questions?

Thank you very much!

www.liu.se

www.liu.se/ouliu

du säger att en människo ej
an forlora vörren efter vad hon
vet eller vad hon
får vet.

Hic sita sum
Quae frugite mi
cum coniguo terrae
has colui, semper mea
dilecta marito.



Lpt 94

Källskriv och rättsteknik
- Attärätt
- Offentlig rätt EG/EU-rätt
- Rättslära, rättslära/övning rätt

