



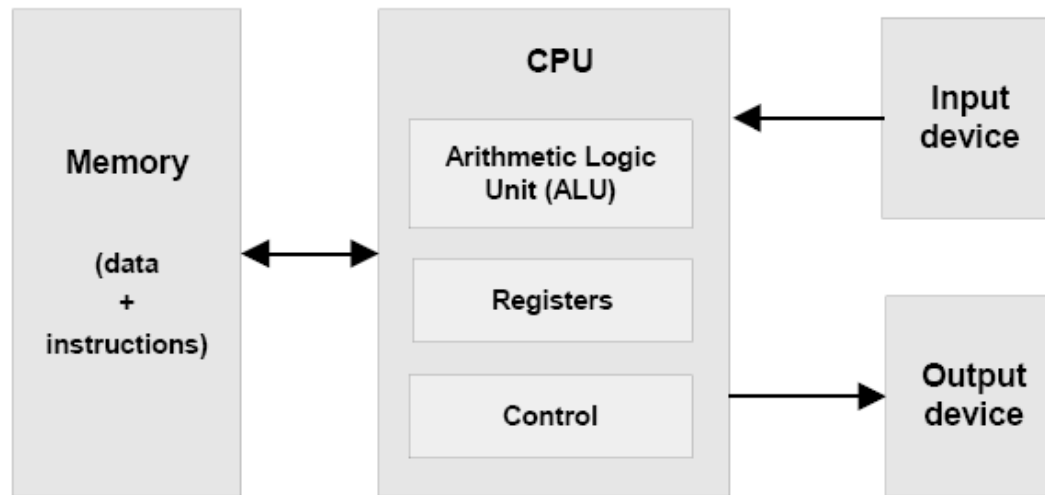
Streaming Processor



Computing Architectures

- General Central Processing Unit (CPU)
- Single Instruction Multiple Data (SIMD)
- Vector processors
- Stream processors


CPU



- Von Neumann processor
- Performance issue
 - ALUs faster than data communication
 - Between memory and ALU
 - Management and communication cost exceeds pure computation cost

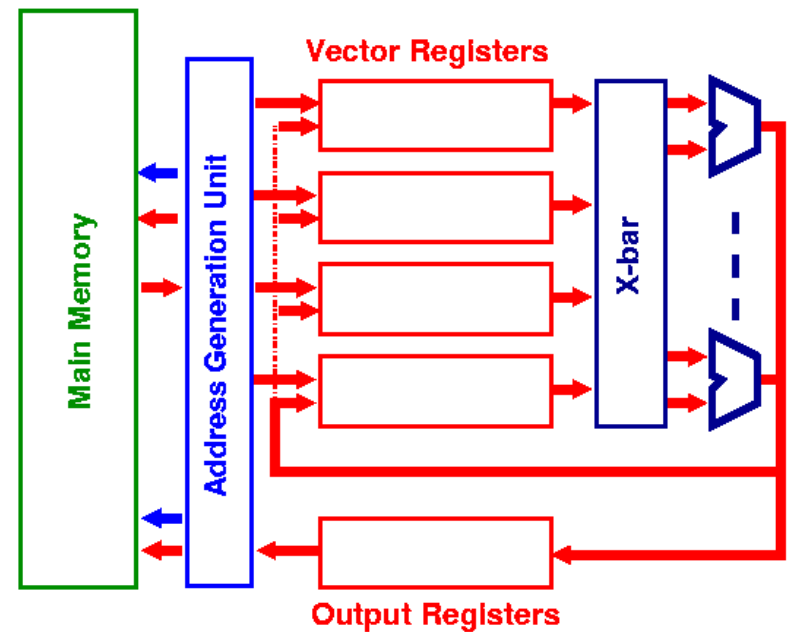
SIMD

- Single Instruction Multiple Data
 - a **single** set of instructions is executed by multiple processors
 - **multiple** data streams are processed at multiple processing units
- Intel MMX, SSE
- AMD 3D Now!
- ...



Vector processors

- Vector data for math computation
- Load vectors with a single instruction
- Operate on multiple data elements simultaneously
- Cray supercomputers





Why Use Stream Processors

- Fast computing by today's VLSI technology
 - thousands of arithmetic logic units operating at multiple GHz on 1cm² die
- Communication and control are bottleneck
 - instructions and data management
- Example:
 - only 6.5% of the Intel Itanium die is devoted to its **12 integer and 2 floating point ALUs** and their registers
 - remainder is used for communication, control, and storage



Why Use Stream Processors


- This is a drawback for general-purpose CPU
- On a GPU, such as a Nvidia GeForce4
 - Hundreds of ALUs
 - Efficient control and communication
 - Special purpose for media-processing (such as graphics pipelines)
 - Exposes abundant parallelism with little global communication and storage



Stream Processors

- Design such processors:
 - Expose these patterns of communication, control, and parallelism to more applications
 - Create a general purpose streaming architecture without compromising its advantages
- Existing implementations that come close
 - Nvidia FX, ATI Radeon GPUs ...
 - enable GP-GPU (general purpose streaming, GP-GPU)

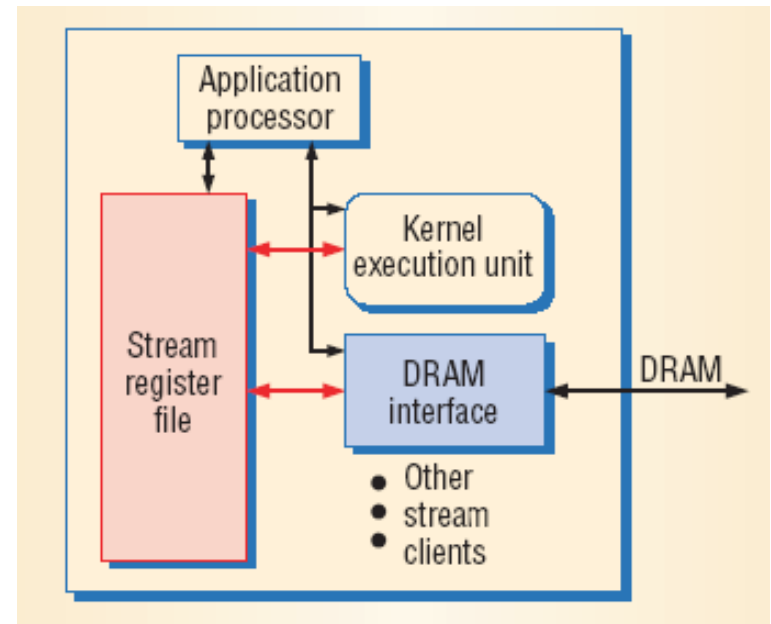
Stream Processing



- Organize an application into streams and kernels
 - inherent locality and concurrency
 - media-processing applications
- This creates the programming model for stream processors

Memory Hierarchy

- Local register files (LRFs)
 - use for operands for arithmetic operations (similar to caches on CPUs)
 - exploit fine-grain locality
- Stream register files (SRFs)
 - capture coarse-grain locality
 - efficiently transfer data to and from the LRFs
- Off-chip memory
 - store global data
 - only use when necessary

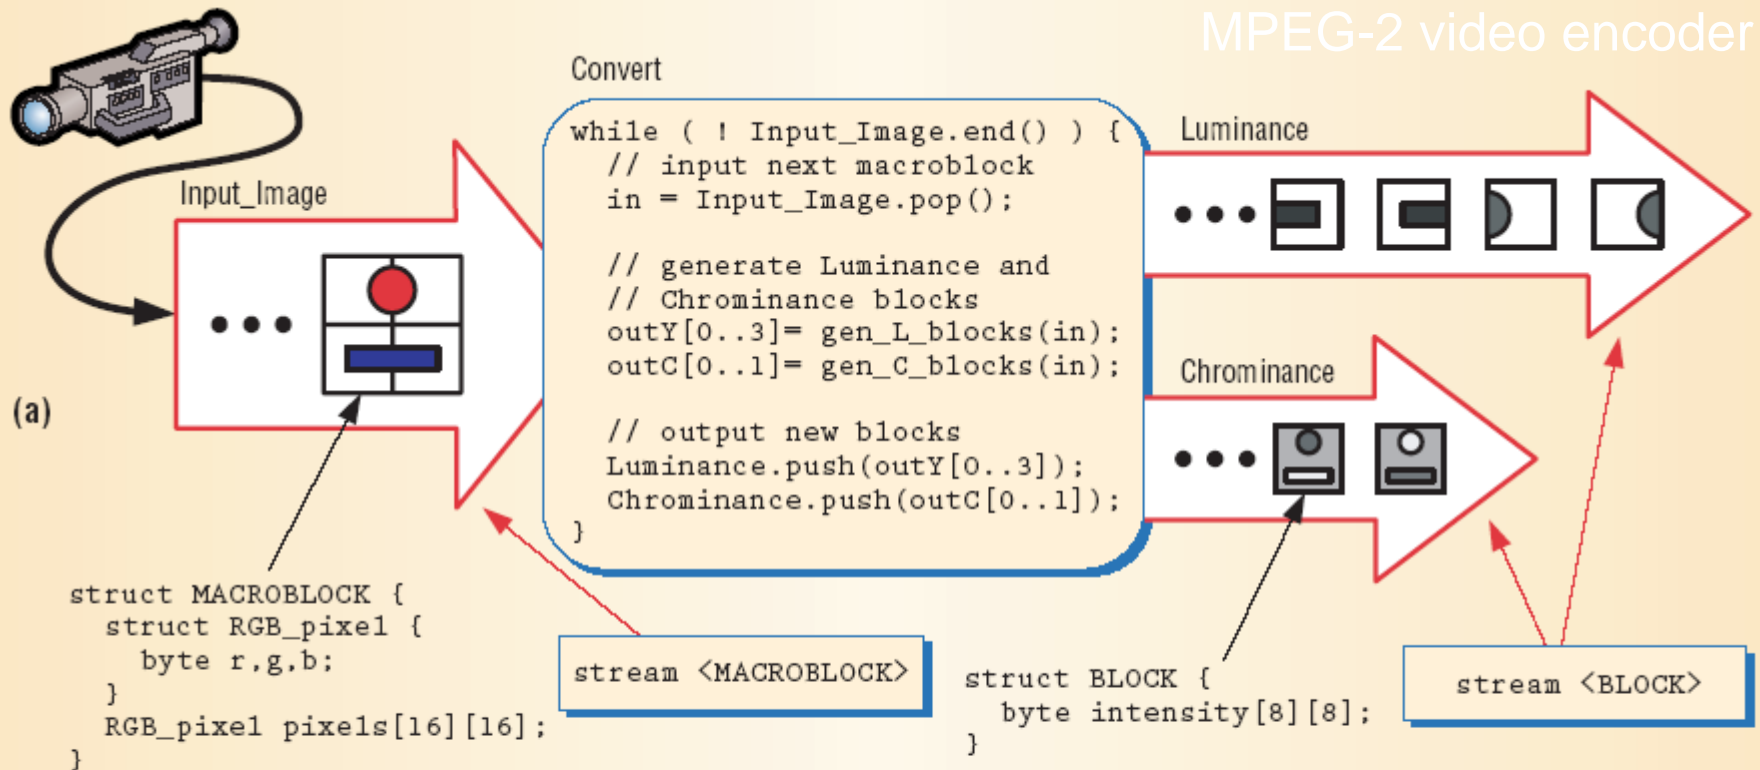




Bandwidth Hierarchy

- Together with the memory hierarchy
 - roughly an order of magnitude for each level
- From today's VLSI technology
- With the locality of operations within LRFs, hundreds of ALUs operate at peak rate

Streams and kernel for MPEG2 encoder



(b)

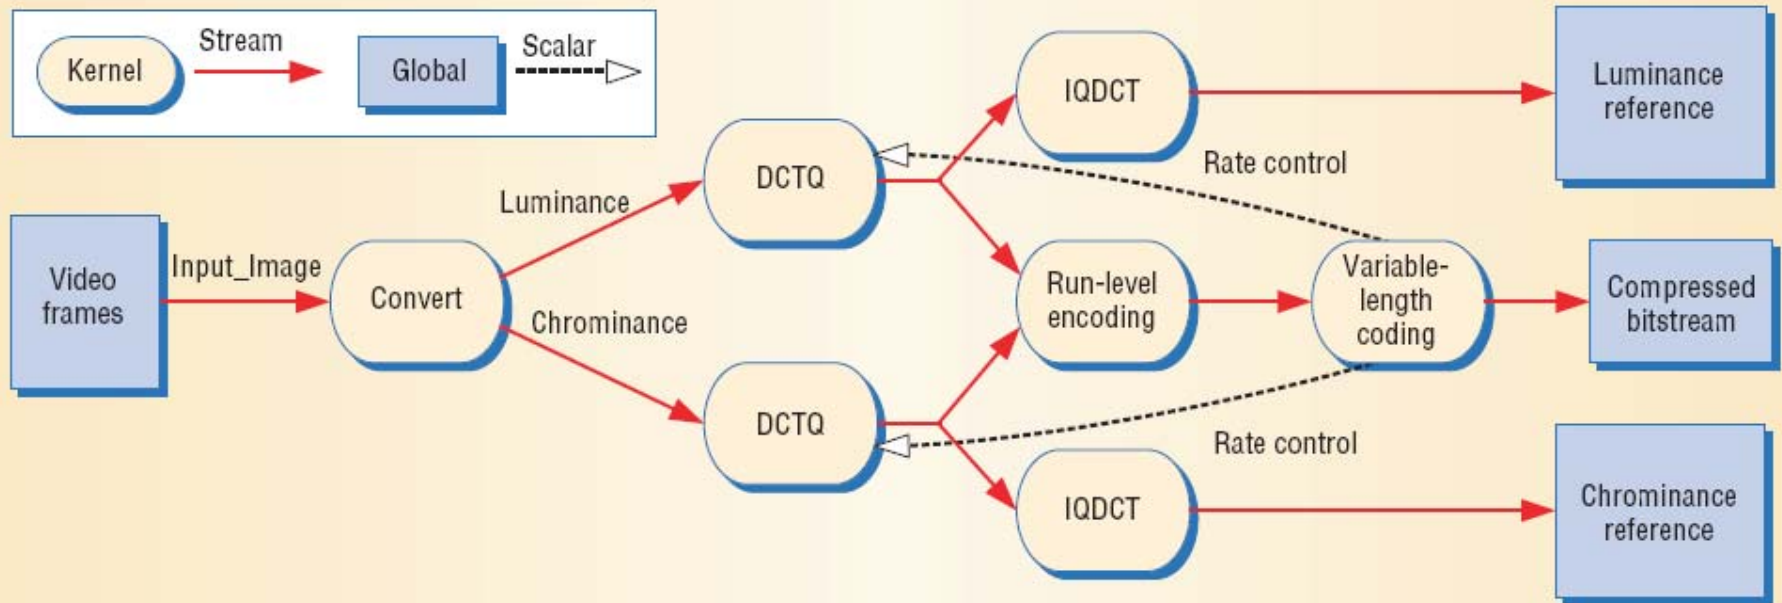
```

stream <MACROBLOCK> Input_Image(NUM_MB);
stream <BLOCK> Luminance(NUM_MB*4), Chrominance(NUM_MB*2),

Input_Image = Video_Feed.get_macroblocks(currpos, NUM_MB);
currpos += NUM_MB;
Convert(Input_Image, Luminance, Chrominance);
    
```

Stream-C program

MPEG2 Encoder

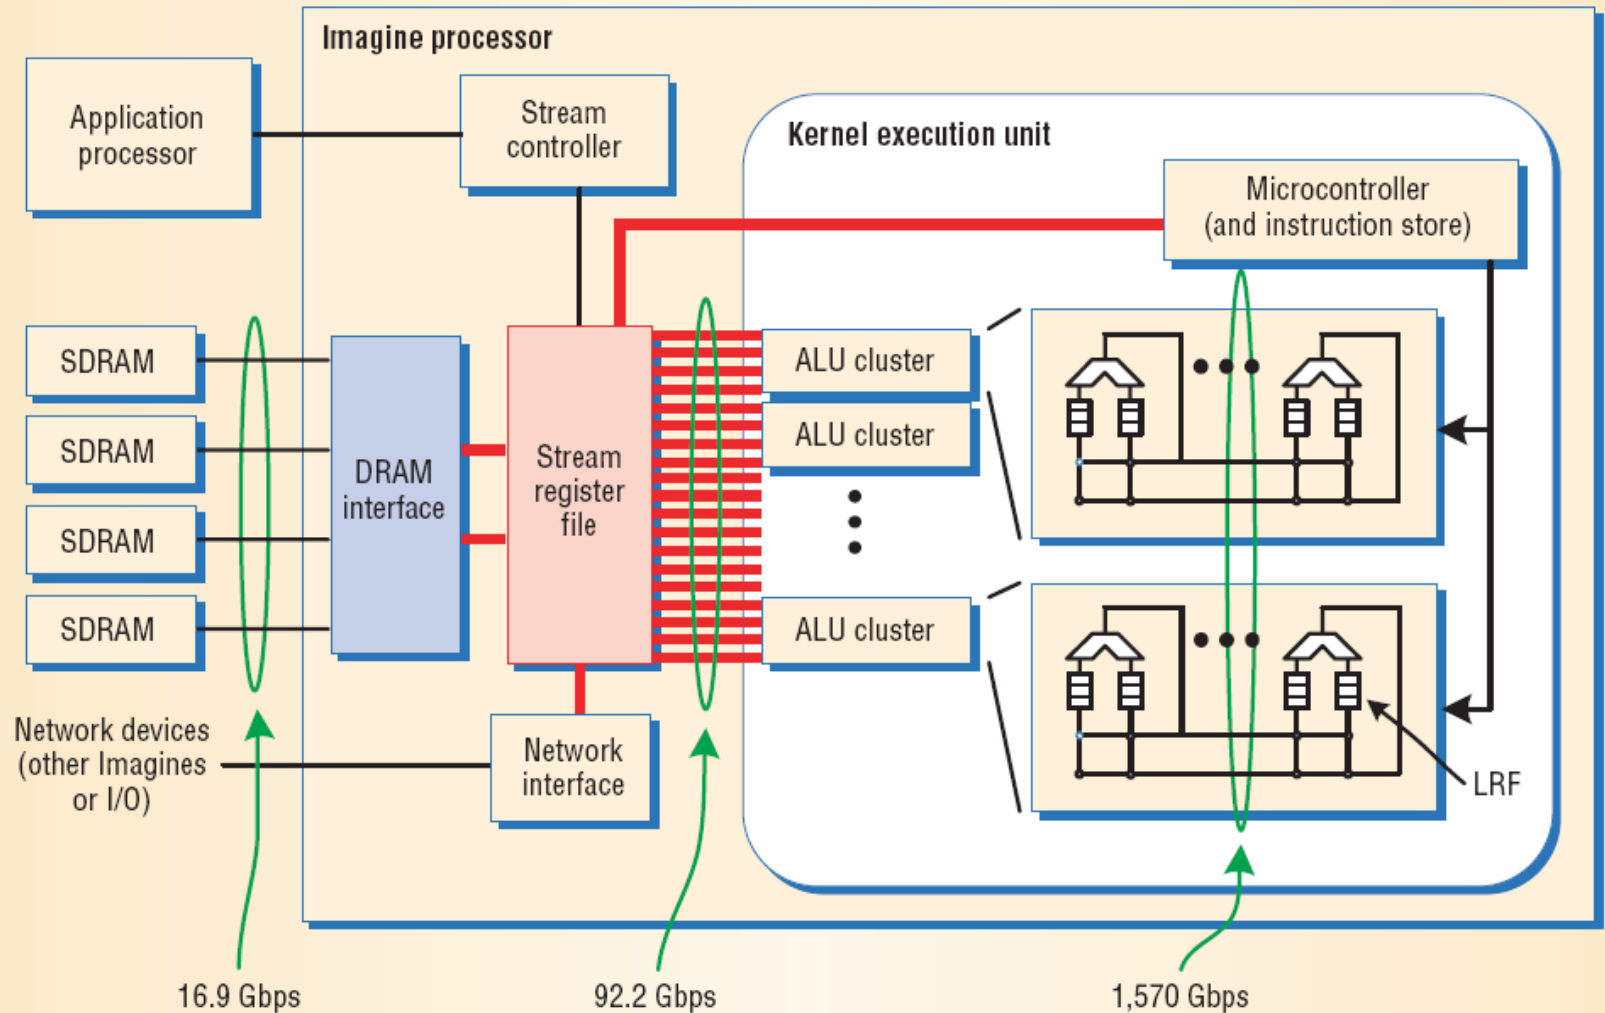


MPEG-2 I-frame encoder

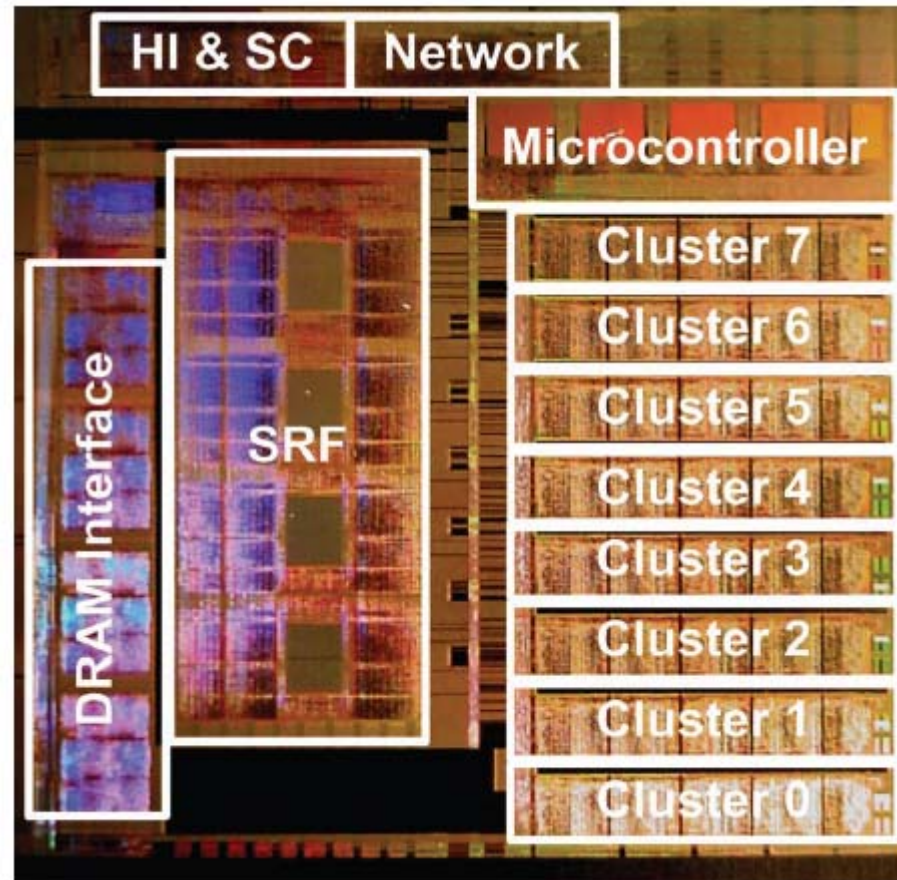
Q: Quantization, IQ: Inverse Quantization, DCT: Discrete Cosine Transform

Global communication (from RAM) needed for the reference frames

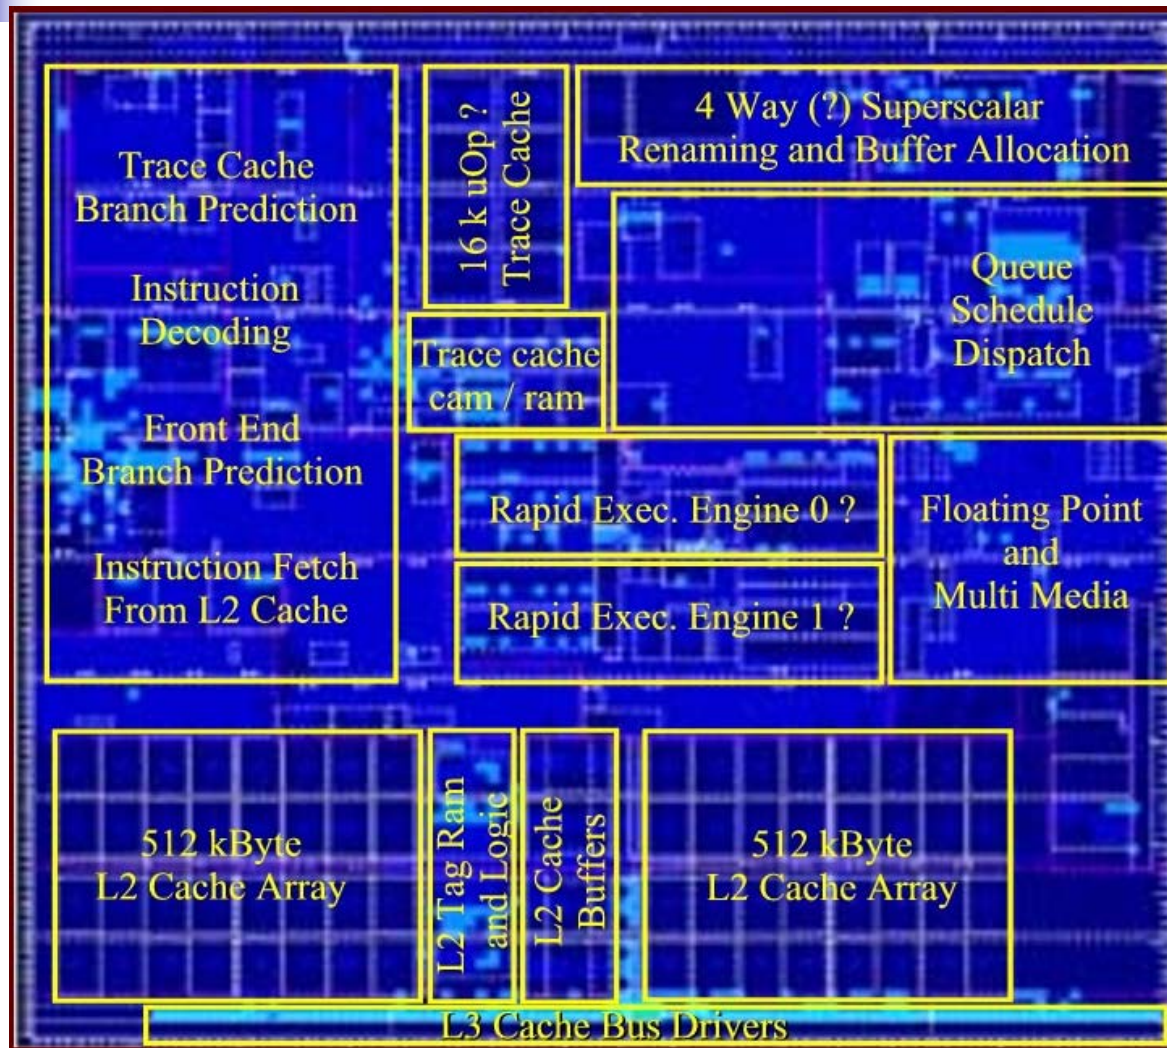
Stanford's Imagine Processor



Imagine's die layout



Intel's Prescott die layout





Parallelism in Stream Processor

- Instruction-level

- exploit parallelism in the scalar operations within a kernel
- for example, gen_L_blocks, gen_C_blocks can occur in parallel

- Data-level

- operate on several data items within a stream in parallel
- for example, different blocks can be converted simultaneously

- Task parallelism

- Multiple tasks runs concurrently
- for example, the two DCTQ kernels could run in parallel



GPUs as Stream Processors

- Stream data elements
 - Points or vertices in vertex processing
 - fragments, essentially pixels in fragment processing
- Kernels
 - vertex and fragment shaders (computing unit)
- Memory
 - texture memory (SRFs)
 - not-exposed LRF
 - bandwidth to RAM (AGP and PCI-Express)



GPUs

- Data parallelism
 - fragments and points are processed in parallel
- Task parallelism
 - fragment and vertex shaders work in parallel
 - data transfer from RAM can be overlapped with computation



References

- U. Kapasi, S. Rixner, W. Dally et al.
“Programmable stream processors,” IEEE Computer August 2003
- S.Venkatasubramanian, “The graphics card as a stream computer,” SIGMOD DIMACS, 2003